Poodle Just-in-time Model Replacement

screenshot

VALUE

SAVINGS AT MAX REQUESTS	SAVINGS %
$627.46	+85.3%

TIMELINE

BREAK-EVEN	EARLIEST SWITCH
87K reqs	2.8 min

TRADE-OFF

	LLM	POODLE
ACCURACY	0.93	0.91
ITEMS / S	13	254

Model Prices

MODEL	INPUT $/1M	OUTPUT $/1M	PROVIDER
BERT 80M small	$0.01	$0.01	Together AI
Llama 3.1 8B large	$0.20	$0.20	Together AI
Llama 3.3 70B Turbo	$0.88	$0.88	Together AI
Llama 3.1 405B Turbo	$3.50	$3.50	Together AI
GPT-4.1	$2.00	$8.00	OpenAI
GPT-4.1 Mini	$0.40	$1.60	OpenAI

Model name | Provider | Input $/1M | Output $/1M | + Add

Models

Large Model: Llama 3.1 8B

Small Model: BERT 80M

Requests

Expected Requests 10,000,000: 100000

Switch After N Requests 5,000: 5000

Use Case

imdb

Input

"A Mouse in the House" is a very classic cartoon by Tom & Jerry, faithful to their tradition but with jokes of its own. It is hysterical, hilarious, very entertaining and quite amusing. Artwork is of good quality either. This short isn't just about Tom trying to catch Jerry. Butch lives in the same house and he's trying to catch the mouse too, because «the's only going to be one cat in this house in the morning - and that's the cat that catches the mouse».

If you ask me, there are lots of funny gags in this cartoon. The funniest for me are, for example, when Mammy Two Shoes sees the two lazy cats

Prompt

You are a sentiment analysis assistant. Your job is to read movie reviews and classify their sentiment as either 'positive' or 'negative'. Only respond in this exact JSON format: {"sentiment": "positive"} or {"sentiment": "negative"}. Do not provide any explanation or

Output

{"sentiment": "positive"}

Model Development

s-500

Model Dev Costs ($) $5.00: 5

Validation

Validation Requests Percent 9%: 0,09

Task Detection

User Provided | Wrapped Requests | 100% | 1

Wrapper Prompt

You are a powerful language model. You are given a user request consisting of a system prompt and a user message. Your task is as follows:

1. Identify the type of input the user is providing: one of ["text", "image", "table", "other"]
2. Infer what task you are expected to perform, choosing from:

Wrapped Output

```
{
  "input_type": "text",
  "task_type": "sentiment classification",
  "user_response": {
    "sentiment": "positive"
  }
}
```